

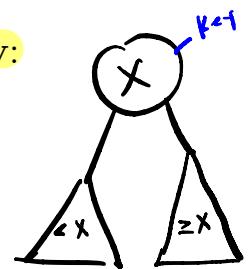
// Balancing Binary Search Trees

Binary Search Trees

- Represented by a linked data structure of nodes.
- $\text{root}(T)$ points to the root of the tree T .
- Each node contains fields:
 - Key
 - left which is the pointer to left child node
(NIL if no left child exists)
 - right which is the pointer to right child node
(NIL if no right child exists)
 - p which is the pointer to parent ($p[\text{root}[T]] = \text{NIL}$)
 - Also whatever satellite data we're storing in the tree

✗ Stored keys must satisfy the **binary search tree property**:

- $\forall y$ in left subtree of x , then $\text{key}[y] < \text{key}[x]$
- $\forall y$ in right subtree of x , then $\text{key}[y] \geq \text{key}[x]$



Find the node in our tree that contains key k using the following algorithm.

Algorithm 1 node TreeSearch(node x , int k)

```
1: if  $x == \text{NIL}$  or  $k == \text{key}[x]$  then
2:   return  $x$ ;
3: end if
4: if  $k < \text{key}[x]$  then
5:   return TreeSearch( $\text{left}[x]$ ,  $k$ );
6: else
7:   return TreeSearch( $\text{right}[x]$ ,  $k$ );
8: end if
```

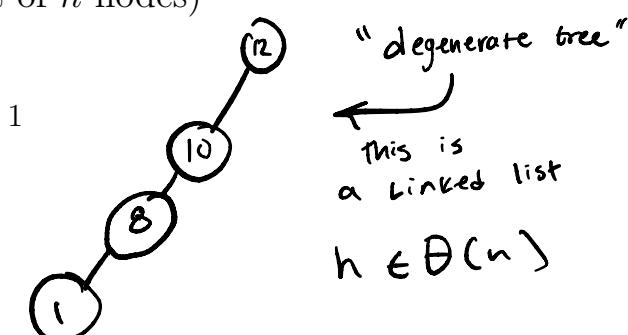
The runtime of the above algorithm depends on the height, h , of the tree T . Specifically, it runs in time $O(h)$.

If the tree is balanced $h \in \Theta(\log n)$

(the distance from the root to any leaf is roughly equal)

In the worst case $h \in \Theta(n)$

(the tree looks like a linked list of n nodes)



// different than hash tables,
as we don't care about #
of keys. Our tree will be built
w/ any # of keys

dynamic size
data structure

We can ensure the tree is balanced using **red-black trees**. This data structure uses an extra one-bit color field. // does not take up much space

worst case

The cost of inserting, deleting, and searching for keys will then take $O(\log n)$ time. // insert / delete are slightly different from normal B.S.T

Red-black tree properties:

(1) Every node is red or black

(2) The root is a black node

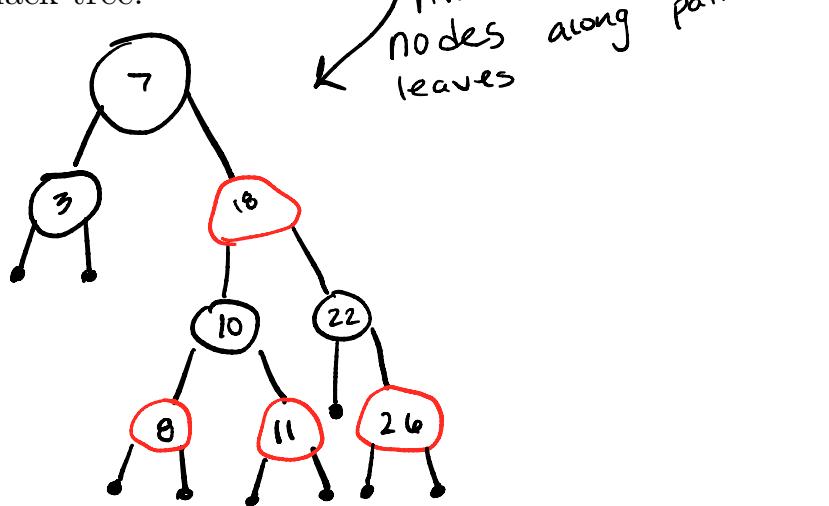
(3) The leaves (NILs) are black nodes

(4) If a node is red, then both of its children are black nodes



(5) All simple paths from the root down to a leaf of the tree contain the same number of black nodes.

Example of a red-black tree:



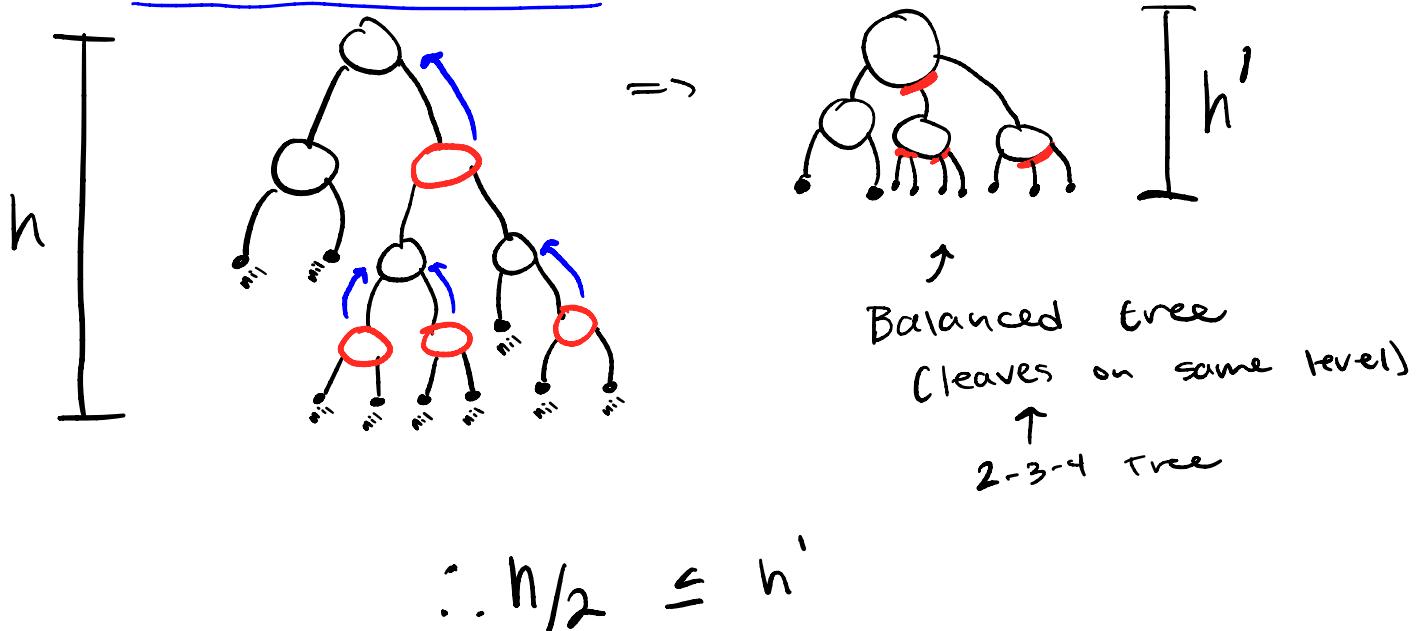
=> Tree does not appear balanced, but because it is a red-black tree - it is.

Theorem.

A red-black tree with n keys has height $h \leq 2 \log(n+1)$.

$$\therefore h \in \Theta(\log n)$$

Proof: Let's imagine collapsing the red nodes of a red-black tree into their black parent nodes



- Collapsing the red nodes yields a tree where every node has 2, 3, or 4 children (i.e a 2-3-4 tree).
- Let h denote the height of the original red-black tree.
- Let h' denote the height of the 2-3-4 tree.
- The leaves of the 2-3-4 tree are all on the same level (from property 5 of red-black trees).
- $h' \geq h/2$ (from property 4 of red-black trees) worst case: every other node is red
- The number of NILs in both trees is $n+1$ (can prove this with an inductive argument). ↳ n is # of keys

$$\Rightarrow n = 8, \# \text{ NILs} = 9$$

$$\Rightarrow n+1 \geq 2^{h'}$$

↳ # NILs if
every node in 2-3-4 tree
has 2 children

$$\Rightarrow \log_2(n+1) \geq h' \geq h/2$$

$$\Rightarrow 2 \log_2(n+1) \geq h \quad \square$$

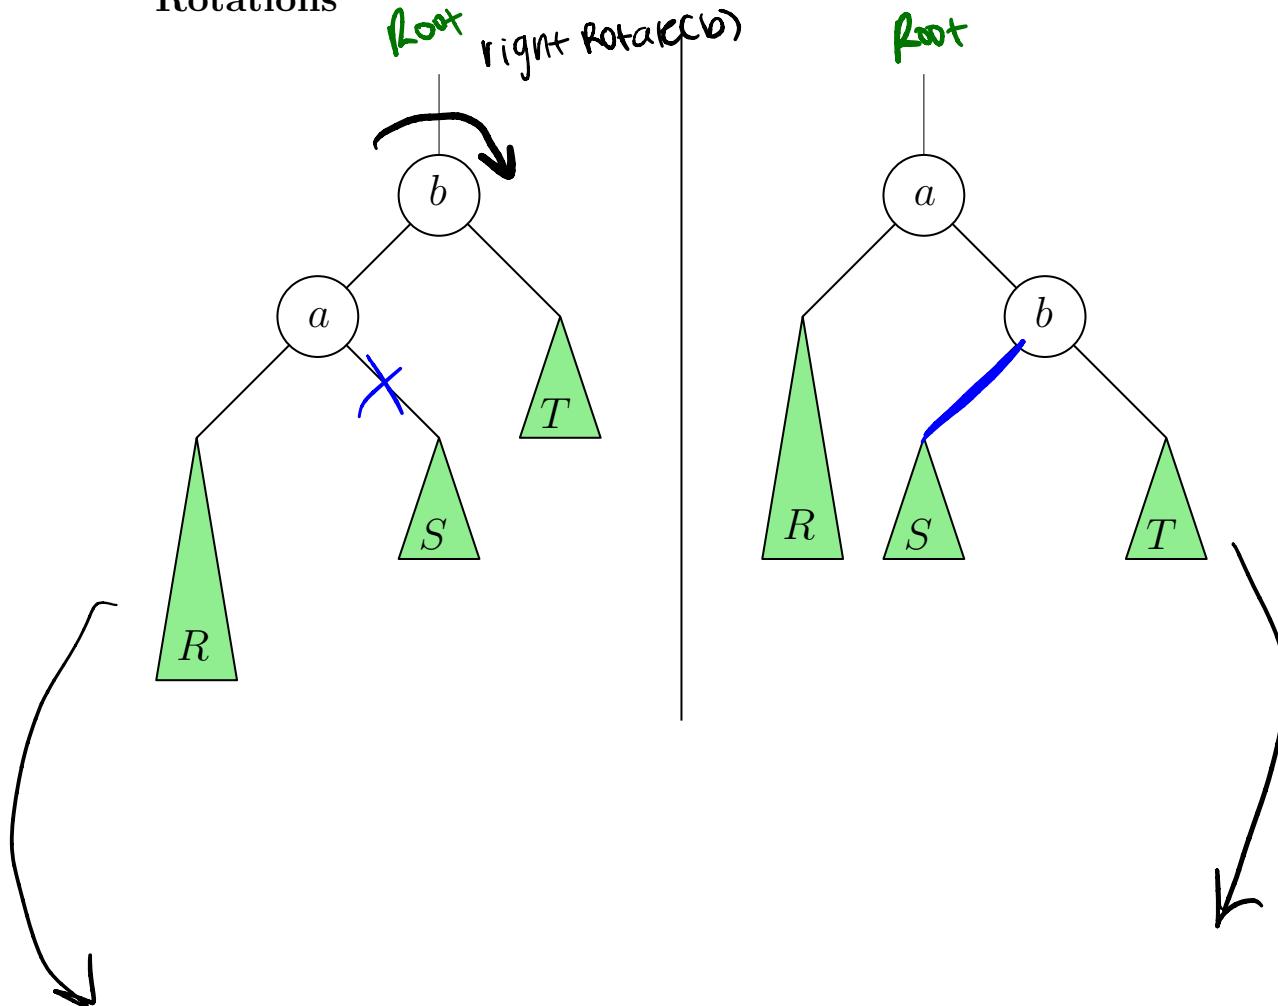
// upper bound of h

The operations Insert and Delete cause modifications to the red-black tree:

- The operation itself.
- Nodes may change colors.
- May require restructuring the links of the tree using rotations.

// R, S, T are large sub trees

Rotations



Rotations maintain the binary search tree property:

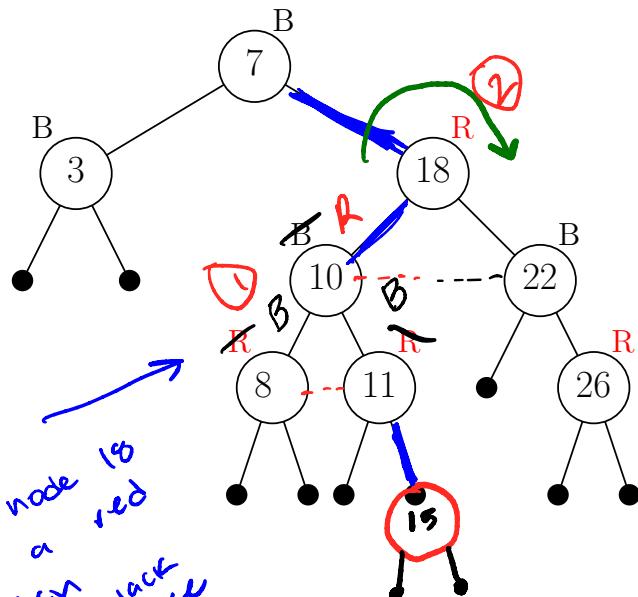
$\forall r \in R, s \in S, t \in T : r < a \leq s < b \leq t$
 \Rightarrow change only 3 pointers to rotate nodes

Rotations can be performed in time $\Theta(1)$

Insertion into a red-black tree

IDEA: Insert x in tree. Color x red. Only red-black property 4 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

Example.

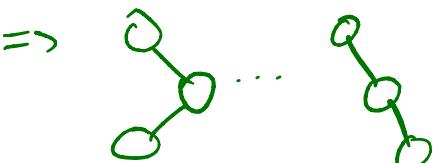


insert (15) ①

$\therefore 15$ placed in that location as it satisfies B.S.T

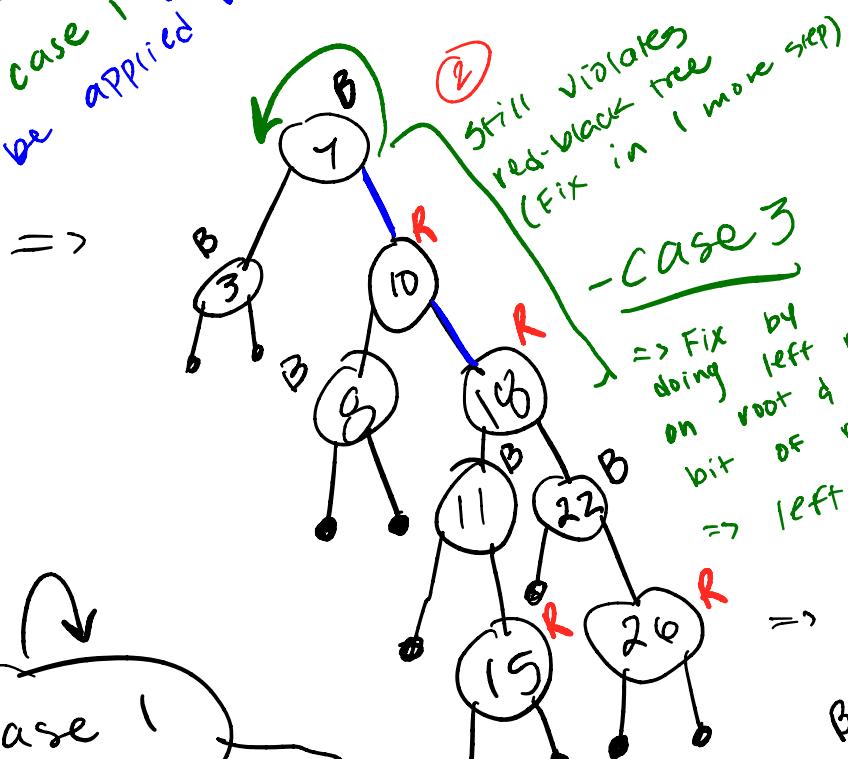
// This violates red-black tree properties ... must fix

\therefore case 2 ②



\Rightarrow right rotate (18)

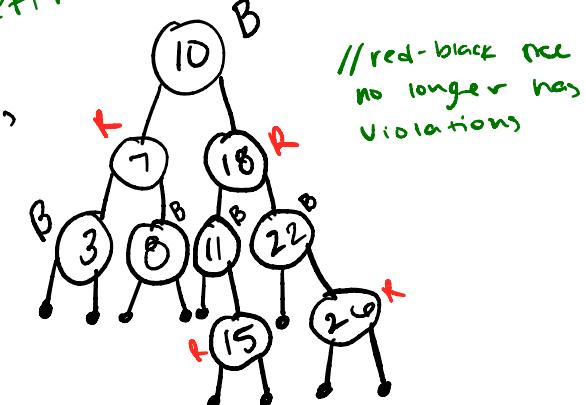
// Not changing colors, instead moving red nodes around



case 1

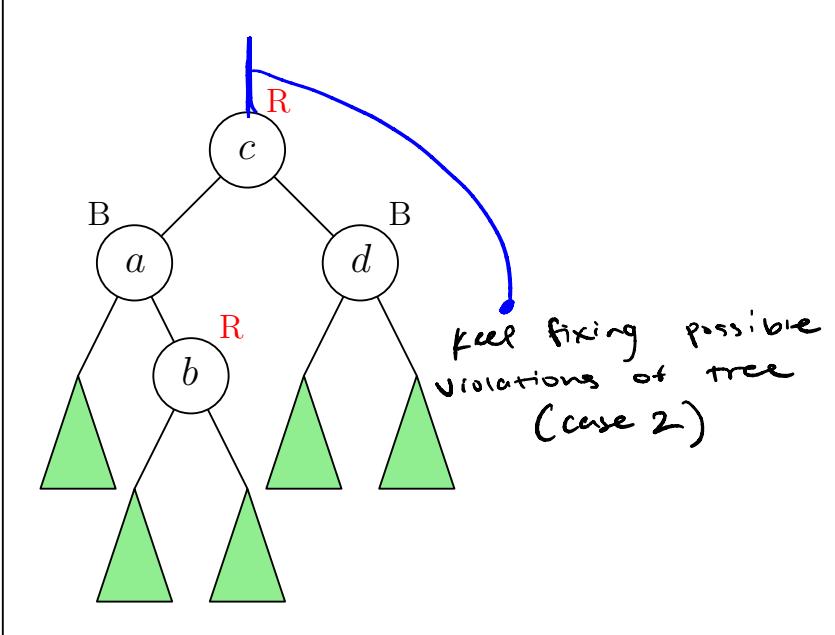
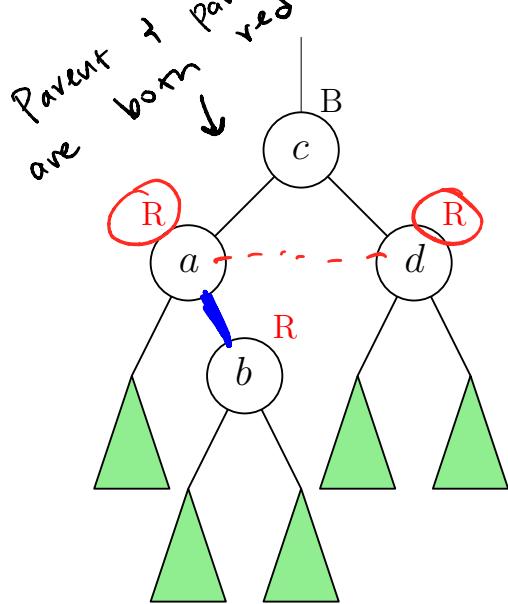
case 4 ⑤

\Rightarrow case 2 \rightarrow case 3 \rightarrow done



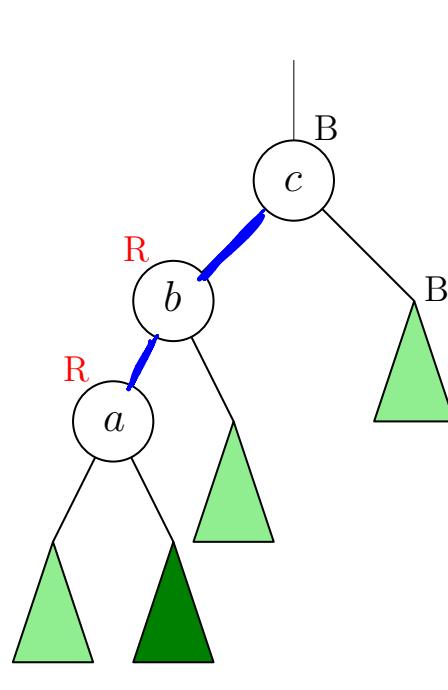
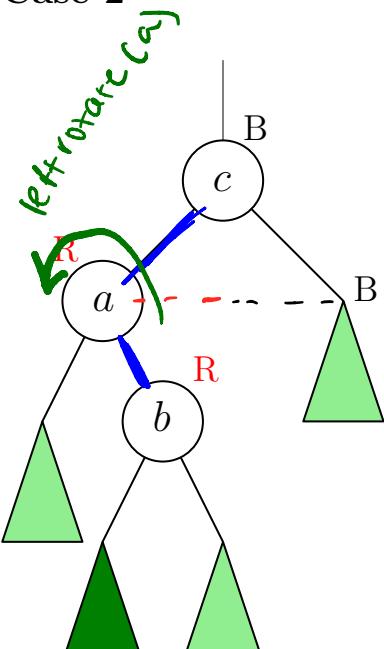
The insertion algorithm considers the following 3 cases (and their left-right mirrors):

Case 1
 Parent & parent's sibling
 are both red



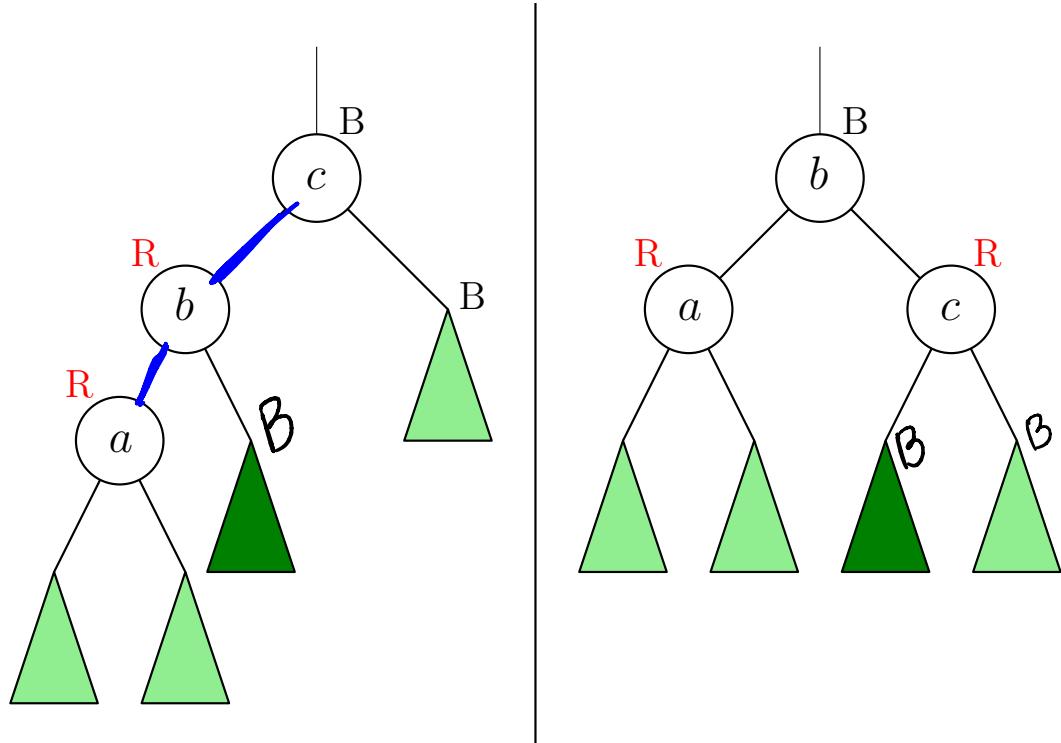
- (continue fixing violations up the root)

Case 2



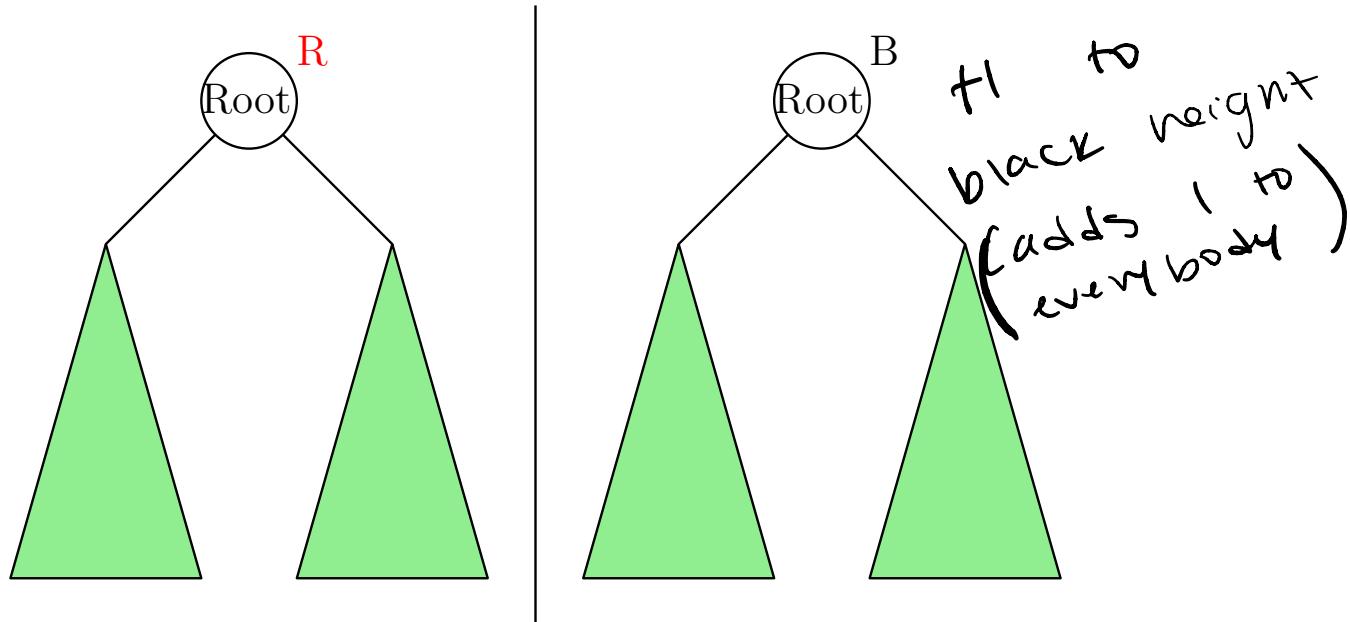
(case 2 always leads directly into case 3)

Case 3



- (no more work is required at this point)

Case 4



(no more work is required at this point)

\Rightarrow If root ever red, just change
 \rightarrow black

// insert(x) $\in \Theta(\log n)$

Analysis

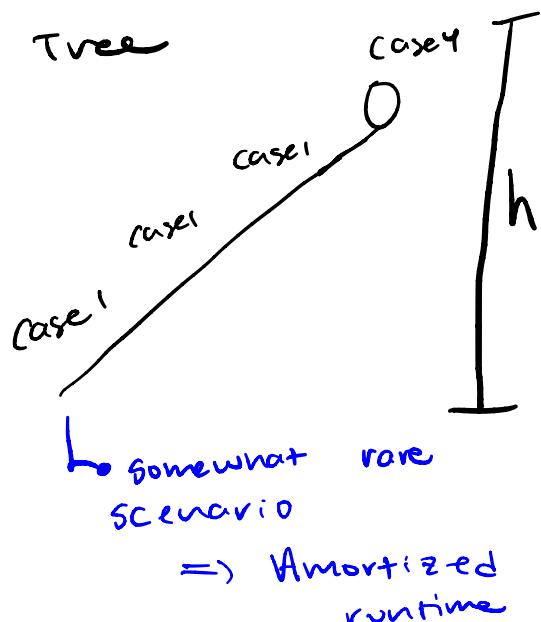
Case 2 and case 3 require only 1 or 2 rotations to fix the tree.

Case 4 only requires only one recolor to fix the tree.

Each case 1 requires three recolor but it doesn't fix the tree and instead pushes it further up the tree.

\Rightarrow Worst case requires $O(\log n)$ operations since we may need to push the problem all the way to the top of the tree.

↳ when you get case 1 all the way up the tree



∴ CIRs for delete(x)
if interested.